

# Assignment2

1. Use Scikit-learn for classification on CIFAR10
2. Use Pytorch to reimplement Lenet for classification on CIFAR10

## Introduction to CIFAR-10

The CIFAR-10 dataset contains 60,000 color images of 32 x 32 pixels in 3 channels divided into 10 classes. Each class contains 6,000 images. The training set contains 50,000 images, while the test sets provides 10,000 images. This image taken from the CIFAR repository (<https://www.cs.toronto.edu/~kriz/cifar.html>). This is a classification problem with 10 classes(multi-label classification). We can take a view on this image for more comprehension of the dataset.

## Use Scikit-learn for classification on CIFAR10

Import library we need

```
import numpy as np
from keras.datasets import cifar10
from keras.utils import to_categorical
import warnings
from sklearn.exceptions import ConvergenceWarning
from sklearn.neural_network import MLPClassifier
import matplotlib.pyplot as plt
```

Load train and test data

```
(X_train, y_train), (X_test, y_test) = cifar10.load_data()

Downloading data from https://www.cs.toronto.edu/~kriz/cifar-10-
python.tar.gz
170498071/170498071 [=====] - 3s 0us/step
```

Look through the shape of the data

```
print("Shape of training data:")
print(X_train.shape)
print(y_train.shape)
print("Shape of test data:")
print(X_test.shape)
print(y_test.shape)
```

```
Shape of training data:
(50000, 32, 32, 3)
(50000, 1)
Shape of test data:
```


```
(10000, 32, 32, 3)
(10000, 1)
```

Visualization of the data `

```
cifar_classes = ['airplane', 'automobile', 'bird', 'cat', 'deer',
                 'dog', 'frog', 'horse', 'ship', 'truck']
print('Example training images and their labels: ' + str([x[0] for x
in y_train[0:5]]))
print('Corresponding classes for the labels: ' +
str([cifar_classes[x[0]] for x in y_train[0:5]]))

f, axarr = plt.subplots(1, 5)
f.set_size_inches(16, 6)
for i in range(5):
    img = X_train[i]
    axarr[i].imshow(img)
plt.show()
```

Example training images and their labels: [6, 9, 9, 4, 1]  
Corresponding classes for the labels: ['frog', 'truck', 'truck',  
'deer', 'automobile']



Normalization

```
# Transform label indices to one-hot encoded vectors
y_train = to_categorical(y_train, num_classes=10)
y_test = to_categorical(y_test, num_classes=10)

# Transform images from (32,32,3) to 3072-dimensional vectors
(32*32*3)
X_train = np.reshape(X_train, (50000, 3072))
X_test = np.reshape(X_test, (10000, 3072))
X_train = X_train.astype('float32')
X_test = X_test.astype('float32')

# Normalization of pixel values (to [0-1] range)
X_train /= 255
X_test /= 255
```

Define model

## MLP Structure

For the MLP network, we build a fully connected network with 2 hidden layers and use relu as the activation function. In the first layer a 32323 image is expanded into a one-dimensional vecto. In the output layer output a 10\*1 matrix and calculate its probability distribution using softmax.

```
from keras.models import Sequential
from keras.layers import Dense, Activation
from tensorflow.keras.optimizers.legacy import SGD

model = Sequential()
model.add(Dense(256, activation='relu', input_dim=3072))
model.add(Dense(256, activation='relu'))
model.add(Dense(10, activation='softmax'))
sgd = SGD(learning_rate=0.01, decay=1e-6, momentum=0.9, nesterov=True)
model.compile(optimizer=sgd, loss='categorical_crossentropy', metrics=['accuracy'])
```

## Train model


```
# Training the MLP
history = model.fit(X_train,y_train, epochs=10, batch_size=32,
verbose=1, validation_split=0.2)

Epoch 1/10
1250/1250 [=====] - 21s 16ms/step - loss:
1.8356 - accuracy: 0.3361 - val_loss: 1.7209 - val_accuracy: 0.3860
Epoch 2/10
1250/1250 [=====] - 14s 11ms/step - loss:
1.6672 - accuracy: 0.4029 - val_loss: 1.6557 - val_accuracy: 0.4124
Epoch 3/10
1250/1250 [=====] - 11s 9ms/step - loss:
1.5956 - accuracy: 0.4297 - val_loss: 1.5814 - val_accuracy: 0.4350
Epoch 4/10
1250/1250 [=====] - 12s 9ms/step - loss:
1.5537 - accuracy: 0.4419 - val_loss: 1.5774 - val_accuracy: 0.4393
Epoch 5/10
1250/1250 [=====] - 12s 9ms/step - loss:
1.5186 - accuracy: 0.4572 - val_loss: 1.6368 - val_accuracy: 0.4139
Epoch 6/10
1250/1250 [=====] - 12s 9ms/step - loss:
1.4874 - accuracy: 0.4665 - val_loss: 1.5803 - val_accuracy: 0.4428
Epoch 7/10
1250/1250 [=====] - 12s 9ms/step - loss:
1.4677 - accuracy: 0.4724 - val_loss: 1.5843 - val_accuracy: 0.4305
Epoch 8/10
1250/1250 [=====] - 12s 9ms/step - loss:
1.4468 - accuracy: 0.4815 - val_loss: 1.5532 - val_accuracy: 0.4490
Epoch 9/10
```

```
1250/1250 [=====] - 12s 10ms/step - loss:
1.4288 - accuracy: 0.4896 - val_loss: 1.5740 - val_accuracy: 0.4439
Epoch 10/10
1250/1250 [=====] - 14s 12ms/step - loss:
1.4099 - accuracy: 0.4978 - val_loss: 1.5131 - val_accuracy: 0.4734
```

Plot model loss

```
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'validation'], loc='upper left')
plt.show()
```



Evaluate Model

```
# Evaluating the MLP
score = model.evaluate(X_test, y_test, batch_size=128, verbose=0)
```

```
print(model.metrics_names)
print(score)

['loss', 'accuracy']
[1.4909449815750122, 0.47679999470710754]
```

## Use Pytorch to reimplement Lenet for classification on CIFAR10

Import library

```
import torch
import torchvision
import torchvision.transforms as transforms
import matplotlib.pyplot as plt
import numpy as np
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
```

Download data

```
# transforms.Compose() splices the two functions together.

transform =
transforms.Compose([transforms.ToTensor(), transforms.Normalize((0.5,
0.5, 0.5), (0.5, 0.5, 0.5))])
trainset = torchvision.datasets.CIFAR10(root='../data',
train=True, download=True, transform=transform)
```

```
trainloader = torch.utils.data.DataLoader(trainset,
batch_size=4, shuffle=True, num_workers=0)
```

```
# Download test data
testset = torchvision.datasets.CIFAR10(root='../data',
train=False, download=False, transform=transform)
testloader = torch.utils.data.DataLoader(testset,
batch_size=4, shuffle=False, num_workers=0)
classes = ('plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog',
'horse', 'ship', 'truck')
```

Downloading <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz> to  
../data/cifar-10-python.tar.gz

100%|██████████| 170498071/170498071 [00:01<00:00, 92672470.46it/s]

Extracting ../data/cifar-10-python.tar.gz to ../data

Visualization of the data

```
def imshow(img):
    # normalize: output = (input-0.5)/0.5
    # unnormalize: input = output*0.5 + 0.5
    img = img / 2 + 0.5      # unnormalize
    npimg = img.numpy()

    plt.imshow(np.transpose(npimg, (1, 2, 0)))
    plt.show()

# Obtained train data randomly
# next() Returns the next item of the iterator.
# next()
dataiter = iter(trainloader)
images, labels = next(dataiter)

# torchvision.utils.make_grid() Splice several pictures into one
picture
imshow(torchvision.utils.make_grid(images))

# Print Label
print(' '.join('%5s' % classes[labels[j]] for j in range(4)))
```




horse horse deer car

Define Model

LeNet Structure

The LeNet-5 architecture consists of two sets of convolutional and average pooling layers, followed by a flattening convolutional layer, then two fully-connected layers and finally a

softmax classifier.



- First Layer: The input for LeNet-5 is a 32×32×3 image which passes through the first convolutional layer with 6 feature maps or filters having size 5×5 and a stride of one. The image dimensions changes from 32x32x3 to 28x28x6.
- Second Layer: Then the LeNet-5 applies average pooling layer or sub-sampling layer with a filter size 2×2 and a stride of two. The resulting image dimensions will be reduced to 14x14x6.
- Third Layer: Next, there is a second convolutional layer with 16 feature maps having size 5×5 and a stride of 1. In this layer, only 10 out of 16 feature maps are connected to 6 feature maps of the previous layer as shown below.
- Fourth Layer: The fourth layer is again an average pooling layer with filter size 2×2 and a stride of 2. This layer is the same as the second layer (S2) except it has 16 feature maps so the output will be reduced to 5x5x16.
- Fifth Layer: The fifth layer is a fully connected convolutional layer with 120 feature maps each of size 1×1. Each of the 120 units in C5 is connected to all the 400 nodes (5x5x16) in the fourth layer.
- Sixth Layer: The sixth layer is a fully connected layer (F6) with 84 units.
- Output Layer: Finally, there is a fully connected softmax output layer  $\hat{y}$  with 10 possible values corresponding to the digits from 0 to 9.

```
import torch
import torch.nn as nn
import torch.nn.functional as F

# Run on GPU if you have one, CPU if you don't.
device = torch.device("cuda:0" if torch.cuda.is_available() else
"cpu")

class LeNet(nn.Module):
    def __init__(self):
        super(LeNet, self).__init__()

        # convolution layer 1: input image depth = 3, output image
        depth = 16, convolution kernel size = 5 * 5, convolution step size =
        1;
        # 16 represents the output dimension and the number of
```

```

convolution cores
    self.conv1 =
nn.Conv2d(in_channels=3,out_channels=16,kernel_size=5,stride=1)

    # pool layer 1: Maximum Pool, area set size = 2 * 2. Pool step
size = 2
    self.pool1 = nn.MaxPool2d(kernel_size=2,stride=2)

    # convolution layer 2
    self.conv2 =
nn.Conv2d(in_channels=16,out_channels=32,kernel_size=5,stride=1)

    # pool layer 2
    self.pool2 = nn.MaxPool2d(kernel_size=2, stride=2)

    # full connected layer 1: input size = 32 * 5 * 5, output size
= 120
    self.fc1 = nn.Linear(32*5*5,120)

    # full connected layer 2:
    self.fc2 = nn.Linear(120,84)

    # full connected layer 3:
    self.fc3 = nn.Linear(84,10)

def forward(self,x):
    x = F.relu(self.conv1(x)) # input(3, 32, 32) output(16, 28,
28)
    x = self.pool1(x) # output(16, 14, 14)
    x = F.relu(self.conv2(x)) # output(32, 10, 10)
    x = self.pool2(x) # output(32, 5, 5)
    x = x.view(-1, 32 * 5 * 5) # output(32*5*5)
    x = F.relu(self.fc1(x)) # output(120)
    x = F.relu(self.fc2(x)) # output(84)
    x = self.fc3(x) # output(10)
    return x

net=LeNet()
net=net.to(device)
# View the network structure
print("Network Structure")
print(net)

Network Structure
LeNet(
  (conv1): Conv2d(3, 16, kernel_size=(5, 5), stride=(1, 1))
  (pool1): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
  (conv2): Conv2d(16, 32, kernel_size=(5, 5), stride=(1, 1))

```



```

(pool2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
(fc1): Linear(in_features=800, out_features=120, bias=True)
(fc2): Linear(in_features=120, out_features=84, bias=True)
(fc3): Linear(in_features=84, out_features=10, bias=True)
)

```

Train model

```

net = LeNet()
net = net.to(device)
loss_function = nn.CrossEntropyLoss() # Cross-entropy loss function

# optimizer = Adam, learning rate = 0.001
optimizer = optim.Adam(net.parameters(), lr=0.001)

for epoch in range(10): # iterate 10 times
    running_loss = 0.0 # Initialize the loss function value loss=0

    for i, data in enumerate(trainloader, start=0):

        # Get the training data
        inputs, labels = data
        inputs, labels = inputs.to(device), labels.to(device) #
Load data and labels to the GPU/CPU

        # Clear the weight parameter gradient
        optimizer.zero_grad()

        # Forward and reverse propagation
        outputs = net(inputs) # Call the above neural network,
forward propagation
        loss = loss_function(outputs, labels) # the loss function
calculates the difference between the convolutional neural network
value and the original value
        loss.backward() # Call pytorch's automatic back
propagation function to automatically generate gradients
        optimizer.step() # Execute the optimizer to propagate the
gradient back to each network

        # Print loss
        running_loss += loss.item()
        if i % 2000 == 1999: # print every 2000 mini-batches
            print('[%d, %5d] loss: %.3f' % (epoch + 1, i + 1,
running_loss / 2000))
            running_loss = 0.0

    print('Finished Training')

```

```
[1, 2000] loss: 1.845
[1, 4000] loss: 1.563
[1, 6000] loss: 1.478
[1, 8000] loss: 1.432
[1, 10000] loss: 1.398
[1, 12000] loss: 1.347
Finished Training
[2, 2000] loss: 1.276
[2, 4000] loss: 1.250
[2, 6000] loss: 1.250
[2, 8000] loss: 1.229
[2, 10000] loss: 1.211
[2, 12000] loss: 1.223
Finished Training
[3, 2000] loss: 1.129
[3, 4000] loss: 1.143
[3, 6000] loss: 1.131
[3, 8000] loss: 1.128
[3, 10000] loss: 1.132
[3, 12000] loss: 1.117
Finished Training
[4, 2000] loss: 1.047
[4, 4000] loss: 1.047
[4, 6000] loss: 1.049
[4, 8000] loss: 1.064
[4, 10000] loss: 1.044
[4, 12000] loss: 1.054
Finished Training
[5, 2000] loss: 0.982
[5, 4000] loss: 0.996
[5, 6000] loss: 0.984
[5, 8000] loss: 1.010
[5, 10000] loss: 1.007
[5, 12000] loss: 0.995
Finished Training
[6, 2000] loss: 0.906
[6, 4000] loss: 0.930
[6, 6000] loss: 0.940
[6, 8000] loss: 0.985
[6, 10000] loss: 0.958
[6, 12000] loss: 0.965
Finished Training
[7, 2000] loss: 0.885
[7, 4000] loss: 0.894
[7, 6000] loss: 0.914
[7, 8000] loss: 0.915
[7, 10000] loss: 0.930
[7, 12000] loss: 0.920
Finished Training
[8, 2000] loss: 0.849
```

```

[8, 4000] loss: 0.863
[8, 6000] loss: 0.903
[8, 8000] loss: 0.881
[8, 10000] loss: 0.890
[8, 12000] loss: 0.897
Finished Training
[9, 2000] loss: 0.807
[9, 4000] loss: 0.846
[9, 6000] loss: 0.873
[9, 8000] loss: 0.850
[9, 10000] loss: 0.869
[9, 12000] loss: 0.856
Finished Training
[10, 2000] loss: 0.797
[10, 4000] loss: 0.810
[10, 6000] loss: 0.818
[10, 8000] loss: 0.844
[10, 10000] loss: 0.843
[10, 12000] loss: 0.850
Finished Training

```

Test model

```

correct = 0
total = 0

# with torch.no_grad() Indicates that it includes content that does
# not need to compute gradients, nor does it need to be propagated
# backward, saving memory
with torch.no_grad():
    for data in testloader:
        images, labels = data
        images, labels = images.to(device), labels.to(device)
        outputs = net(images)
        # torch.max(outputs.data, 1)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

print('Accuracy of the network on the 10000 test images: %d %%' % (100
* correct / total))

# Print the accuracy of 10 classes
class_correct = list(0. for i in range(10)) #class_correct=[0.0, 0.0,
0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0]
class_total = list(0. for i in range(10)) #class_total=[0.0, 0.0, 0.0,
0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0]
with torch.no_grad():

```

```

for data in testloader:
    images, labels = data
    images, labels = images.to(device), labels.to(device)
    outputs = net(images) # outputs vector: 4*10
    # torch.max(outputs.data,
    # predicted vector: 4*1
    _, predicted = torch.max(outputs, 1)

    c = (predicted == labels).squeeze()
    for i in range(4):
        label = labels[i]
        class_correct[label] += c[i].item()
        class_total[label] += 1
for i in range(10):
    print('Accuracy of %5s : %2d %%' % (classes[i], 100 *
class_correct[i] / class_total[i]))

Accuracy of the network on the 10000 test images: 63 %

```

Save the weights and parameters

```

save_path = 'Lenets.pth'
torch.save(net.state_dict(), save_path)

```

Make predict by the model we train

```

import torch
import torchvision.transforms as transforms
from PIL import Image
from torch import nn
import torch.nn.functional as F

transform = transforms.Compose([transforms.Resize((32,
32)), transforms.ToTensor(), transforms.Normalize((0.5, 0.5, 0.5), (0.5,
0.5, 0.5))])

classes = ('plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog',
'horse', 'ship', 'truck')

class LeNet(nn.Module):
    def __init__(self):
        super(LeNet, self).__init__()
        # Convolution layer 1
        self.conv1 =
nn.Conv2d(in_channels=3, out_channels=16, kernel_size=5, stride=1)
        # Pool Layer 1
        self.pool1 = nn.MaxPool2d(kernel_size=2, stride=2)
        # Convolution layer 2

```

```

        self.conv2 =
nn.Conv2d(in_channels=16,out_channels=32,kernel_size=5,stride=1)
        # Pool layer 2
        self.pool2 = nn.MaxPool2d(kernel_size=2, stride=2)
        # Full connected layer 1
        self.fc1 = nn.Linear(32*5*5,120)
        # Full connected layer 2
        self.fc2 = nn.Linear(120,84)
        # Full connected layer 3
        self.fc3 = nn.Linear(84,10)

    def forward(self,x): # Link the input layer, the network layer and
the output layer for forward data transmission
        x = F.relu(self.conv1(x)) # input(3, 32, 32) output(16, 28,
28)
        x = self.pool1(x) # output(16, 14, 14)
        x = F.relu(self.conv2(x)) # output(32, 10, 10)
        x = self.pool2(x) # output(32, 5, 5)
        x = x.view(-1, 32 * 5 * 5) # output(32*5*5)
        x = F.relu(self.fc1(x)) # output(120)
        x = F.relu(self.fc2(x)) # output(84)
        x = self.fc3(x) # output(10)
        return x

net = LeNet()
net.load_state_dict(torch.load('Lenets.pth'))

im = Image.open('cat.jpg')
im = transform(im) # [C, H, W]

# [batch, channel, height, width],
im = torch.unsqueeze(im, dim=0) # [N, C, H, W]

with torch.no_grad():
    outputs = net(im)
    predict = torch.max(outputs, dim=1)[1].data.numpy()
print(classes[int(predict)])

# Predit result
with torch.no_grad():
    outputs = net(im)
    predict = torch.softmax(outputs,dim=1) #
[batch, channel, height, width],
print(predict)

cat
tensor([[0.0062, 0.0065, 0.0446, 0.3172, 0.0830, 0.2256, 0.0430,
0.1104, 0.1509,
0.0127]])

```